

Marketing Campaigns

Problem scenario:

Marketing mix stands as a widely utilized concept in the execution of marketing strategies. It encompasses various facets within a comprehensive marketing plan, with a central focus on the four Ps of marketing: product, price, place, and promotion.

Problem objective:

As a data scientist, you must conduct exploratory data analysis and hypothesis testing to enhance your comprehension of the diverse factors influencing customer acquisition.

Data description:

The variables such as birth year, education, income, and others pertain to the first 'P' or 'People' in the tabular data presented to the user. The expenditures on items like wine, fruits, and gold, are associated with 'Product'. Information relevant to sales channels, such as websites and stores, is connected to 'Place', and the fields discussing promotions and the outcomes of various campaigns are linked to 'Promotion'.

Step 1: Import Libraries & Load the Dataset

- * Imported essential libraries for data manipulation and visualization
- * Loaded the dataset (marketing_data.csv)
- * Displayed the first few rows to verify successful loading
- * Checked the dataset shape and structure

```
[63]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import pearsonr, spearmanr, mannwhitneyu, pointbiserialr

[20]: # Load dataset
df = pd.read_csv("marketing_data.csv")
df.head()

[20]:   ID Year_Birth Education Marital_Status Income Kidhome Teenhome Dt_Customer Recency MntWines ... NumStorePurchases NumWebVisitsMonth AcceptedCmp3 AcceptedCmp4 AcceptedCmp5
0  1826    1970  Graduation      Divorced $84,835.00     0     0  6/16/14       0    189 ...           6          1        0        0        0
1     1    1961  Graduation       Single $57,091.00     0     0  6/15/14       0    464 ...           7          5        0        0        0
2  10476    1958  Graduation      Married $67,267.00     0     1  5/13/14       0    134 ...           5          2        0        0        0
3  1386    1967  Graduation     Together $32,474.00     1     1  5/11/14       0     10 ...           2          7        0        0        0
4   5371    1989  Graduation       Single $21,474.00     1     0  4/8/14       0      6 ...           2          7        1        0        0

5 rows × 28 columns
```

```
[21]: # Check shape and info
print("Shape:", df.shape)
df.info()

Shape: (2240, 28)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2240 entries, 0 to 2239
Data columns (total 28 columns):
 #   Column            Non-Null Count  Dtype  
 --- 
 0   ID                2240 non-null    int64  
 1   Year_Birth         2240 non-null    int64  
 2   Education          2240 non-null    object  
 3   Marital_Status     2240 non-null    object  
 4   Income             2216 non-null    int64  
 5   Kidhome            2240 non-null    int64  
 6   Teenhome            2240 non-null    int64  
 7   Dt_Customer         2240 non-null    object  
 8   Recency             2240 non-null    int64  
 9   MntLlines            2240 non-null    int64  
 10  MntFruits            2240 non-null    int64  
 11  MntHealthProducts    2240 non-null    int64  
 12  MntFishProducts      2240 non-null    int64  
 13  MntSweetProducts     2240 non-null    int64  
 14  MntGoldProd          2240 non-null    int64  
 15  NumDealsPurchases    2240 non-null    int64  
 16  NumLebPurchases      2240 non-null    int64  
 17  NumCatalogPurchases  2240 non-null    int64  
 18  NumStorePurchases     2240 non-null    int64  
 19  NumLebVisitsMnths     2240 non-null    int64  
 20  AcceptedCmp3          2240 non-null    int64  
 21  AcceptedCmp4          2240 non-null    int64  
 22  AcceptedCmp5          2240 non-null    int64  
 23  AcceptedCmp1          2240 non-null    int64  
 24  AcceptedCmp2          2240 non-null    int64  
 25  Response              2240 non-null    int64  
 26  Complain              2240 non-null    int64  
 27  Country               2240 non-null    object 
```

```
[21]: # Check shape and info
print("Shape:", df.shape)
df.info()

Shape: (2240, 28)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2240 entries, 0 to 2239
Data columns (total 28 columns):
 #   Column            Non-Null Count  Dtype  
 --- 
 0   ID                2240 non-null    int64  
 1   Year_Birth         2240 non-null    int64  
 2   Education          2240 non-null    object  
 3   Marital_Status     2240 non-null    object  
 4   Income             2216 non-null    int64  
 5   Kidhome            2240 non-null    int64  
 6   Teenhome            2240 non-null    int64  
 7   Dt_Customer         2240 non-null    object  
 8   Recency             2240 non-null    int64  
 9   MntLlines            2240 non-null    int64  
 10  MntFruits            2240 non-null    int64  
 11  MntHealthProducts    2240 non-null    int64  
 12  MntFishProducts      2240 non-null    int64  
 13  MntSweetProducts     2240 non-null    int64  
 14  MntGoldProd          2240 non-null    int64  
 15  NumDealsPurchases    2240 non-null    int64  
 16  NumLebPurchases      2240 non-null    int64  
 17  NumCatalogPurchases  2240 non-null    int64  
 18  NumStorePurchases     2240 non-null    int64  
 19  NumLebVisitsMnths     2240 non-null    int64  
 20  AcceptedCmp3          2240 non-null    int64  
 21  AcceptedCmp4          2240 non-null    int64  
 22  AcceptedCmp5          2240 non-null    int64  
 23  AcceptedCmp1          2240 non-null    int64  
 24  AcceptedCmp2          2240 non-null    int64  
 25  Response              2240 non-null    int64  
 26  Complain              2240 non-null    int64  
 27  Country               2240 non-null    object 

dtypes: int64(23), object(5)
memory usage: 490.1+ KB
```

Step 2: Initial Data Inspection

- * Checked column names and identified any irregular formatting
 - * Reviewed summary statistics for numerical variables
 - * Inspected missing values across all columns
 - * Verified that key variables (Income, Dt_Customer) require cleaning

```
[22]: # View column names
df.columns.tolist()

[22]: ['ID',
       'Year_Birth',
       'Education',
       'Marital_Status',
       'Income',
       'Kidhome',
       'Teenhome',
       'Dt_Customer',
       'Recency',
       'MntWines',
       'MntFruits',
       'MntMeatProducts',
       'MntFishProducts',
       'MntSweetProducts',
       'MntGoldProducts',
       'NumDealsPurchases',
       'NumWebPurchases',
       'NumCatalogPurchases',
       'NumStorePurchases',
       'NumVisitsMonth',
       'AcceptedCmp1',
       'AcceptedCmp2',
       'AcceptedCmp3',
       'AcceptedCmp4',
       'AcceptedCmp5',
       'AcceptedCmp6',
       'AcceptedCmp7',
       'Response',
       'Complain',
       'Country']

[23]: # Summary statistics for numeric columns

df.describe()

[23]:
```

[23]: # Summary statistics for numeric columns

```
df.describe()
```

[23]:

	ID	Year_Birth	Kidhome	Teenhome	Recency	MntWines	MntFruits	MntMeatProducts	MntFishProducts	MntSweetProducts	...	NumCatalogPurchases	NumStorePurchases	NumWe
count	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	...	2240.000000	2240.000000	
mean	5592.159821	1968.805804	0.444196	0.506250	49.109375	303.935714	26.302232	166.950000	37.525446	27.062946	...	2.662054	5.790179	
std	3246.662198	11.984069	0.538398	0.544538	28.962453	336.597393	39.773434	225.715373	54.628979	41.280498	...	2.923101	3.250958	
min	0.000000	1893.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	...	0.000000	0.000000	
25%	2828.250000	1959.000000	0.000000	0.000000	24.000000	23.750000	1.000000	16.000000	3.000000	1.000000	...	0.000000	3.000000	
50%	5458.500000	1970.000000	0.000000	0.000000	49.000000	173.500000	8.000000	67.000000	12.000000	8.000000	...	2.000000	5.000000	
75%	8427.750000	1977.000000	1.000000	1.000000	74.000000	504.250000	33.000000	232.000000	50.000000	33.000000	...	4.000000	8.000000	
max	11191.000000	1996.000000	2.000000	2.000000	99.000000	1493.000000	199.000000	1725.000000	259.000000	263.000000	...	28.000000	13.000000	

8 rows × 23 columns

[24]: # Check missing values in each column

```
df.isna().sum().sort_values(ascending=False)
```

[24]:

Income	24
ID	0
NumDealsPurchases	0
Complain	0
Response	0
AcceptedCmp2	0
AcceptedCmp1	0
AcceptedCmp5	0
AcceptedCmp4	0
AcceptedCmp3	0
NumWebVisitsMonth	0
NumStorePurchases	0
NumCatalogPurchases	0
NumWebPurchases	0
MntTotalPrcds	0

[25]: # Preview unique values for key categorical columns

```
print("Education:", df['Education'].unique())
print("Marital_Status:", df['Marital_Status'].unique())
print("Country:", df['Country'].unique())
```

Education: ['Graduation' 'PHD' '2n Cycle' 'Master' 'Basic']
Marital_Status: ['Divorced' 'Single' 'Married' 'Together' 'Widow' 'YOLO' 'Alone' 'Absurd']
Country: ['SP' 'CA' 'US' 'AUS' 'GER' 'IND' 'SA' 'MEX']

Step 3: Data Cleaning

3.1 - Clean Column Names

- * Cleaned column names by removing leading/trailing whitespace
- * Standardized column names by replacing spaces with underscores
- * Ensured consistent naming for later processing

```
[26]: # Clean Column Names
df.columns = df.columns.str.strip().str.replace(' ', '_')
df.columns.tolist()

[26]: ['ID',
       'Year_Birth',
       'Education',
       'Marital_Status',
       'Income',
       'Kidhome',
       'Teenhome',
       'Dt_Customer',
       'Recency',
       'MntWines',
       'MntFruits',
       'MntMeatProducts',
       'MntFishProducts',
       'MntSweetProducts',
       'MntGoldProducts',
       'NumDealsPurchases',
       'NumWebPurchases',
       'NumCatalogPurchases',
       'NumStorePurchases',
       'NumWebVisitsMonth',
       'Acceptedmps',
       'Acceptedmpsa',
       'Acceptedmps1',
       'Acceptedmpm1',
       'Acceptedmp2',
       'Response',
       'Complain',
       'Country']
```

3.2 - Fix Income Formatting

- * Remove \$, , and whitespace
- * Convert Income to numeric
- * Confirm missing-value count

```
[28]: # Fix Income formatting
df['Income'] = (
    df['Income']
    .astype(str)
    .str.replace(r'[\$,]', '', regex=True)
    .str.strip()
)

df['Income'] = pd.to_numeric(df['Income'], errors='coerce')

df['Income'].isna().sum()

[28]: np.int64(24)
```

3.3 - Clean Marital_Status Categories

- * Standardize inconsistent categories (Alone, YOLO, Absurd, etc.)
- * Map them into consistent groups: Married, Single, Previously_Married
- * Ensure categories match the business logic

```
[29]: # Clean Marital_Status categories
marital_map = {
    'Married': 'Married',
    'Together': 'Married',
    'Single': 'Single',
    'Alone': 'Single',
    'YOLO': 'Single',
    'Absurd': 'Single',
    'Divorced': 'Previously_Married',
    'Widow': 'Previously_Married'
}

df['Marital_Status'] = df['Marital_Status'].map(marital_map)

df['Marital_Status'].value_counts()

[29]: Marital_Status
      Married          1444
      Single           487
      Previously_Married   309
      Name: count, dtype: int64
```

3.4 - Clean Education Categories

The dataset sometimes contains inconsistent education labels.

This step will:

- * Standardize education categories
- * Ensure categories match the expected five groups
- * Verify final unique values

```
[32]: # Clean Education categories
edu_map = {
    'Graduation': 'Graduation',
    'PHD': 'PHD',
    '2n Cycle': '2n Cycle',
    'Master': 'Master',
    'Basic': 'Basic'
}

df['Education'] = df['Education'].map(edu_map)

df['Education'].value_counts()
```

Education	count
Graduation	1127
PHD	486
Master	370
2n Cycle	203
Basic	54

3.5 - Convert Dt_Customer to Datetime Format

- * Convert the Dt_Customer field from string to datetime
- * Ensure proper recognition of month/day/year format
- * Verify successful conversion

```
[33]: df['Dt_Customer'] = pd.to_datetime(df['Dt_Customer'], format='%m/%d/%y')

df['Dt_Customer'].head()
```

	Dt_Customer
0	2014-06-16
1	2014-06-15
2	2014-05-13
3	2014-05-11
4	2014-04-08

3.6 - Impute Missing Income Values

According to the problem statement:

- `Customers with similar education and marital status tend to have similar yearly incomes on average.'

So we will:

- * Group by Education + Marital_Status
- * Compute median income for each group
- * Fill missing Income values accordingly
- * Verify no missing values remain

```
[35]: # compute median income by Education + Marital_Status
group_medians = df.groupby(['Education', 'Marital_Status'])['Income'].median()

# fill missing income
df['Income'] = df.apply(
    lambda row: group_medians[row['Education'], row['Marital_Status']],
    if pd.isna(row['Income']) else row['Income'],
    axis=1
)

df['Income'].isna().sum()
```

	Income
0	0

Step 4: Feature Engineering

Create feature columns

- * Derived TotalChildren
- * Calculated Age
- * Created TotalSpending
- * Derived TotalPurchases for all channels

```
[36]: # total children
df['TotalChildren'] = df['Kidhome'] + df['Teenhome']

# age (using 2025 as reference year)
df['Age'] = 2025 - df['Year_Birth']

# total spending across product categories
spending_cols = [
    'MntLanes', 'MntFruits', 'MntMeatProducts',
    'MntFishProducts', 'MntSweetProducts', 'MntGoldProds'
]
df['TotalSpending'] = df[spending_cols].sum(axis=1)

# total purchases across all channels
purchase_cols = ['NumWebPurchases', 'NumCatalogPurchases', 'NumStorePurchases']
df['TotalPurchases'] = df[purchase_cols].sum(axis=1)

df[['TotalChildren', 'Age', 'TotalSpending', 'TotalPurchases']].head()
```

	TotalChildren	Age	TotalSpending	TotalPurchases
0	0	55	1190	14
1	0	64	577	17
2	1	67	251	10
3	2	58	11	3
4	1	36	91	6

Step 5: Outlier Treatment for Income

- * Calculated 1st and 99th percentiles
- * Capped values outside this range
- * Replaced original Income with capped values
- * Verified updated distribution

```
[37]: # compute percentiles
Q1 = df['Income'].quantile(0.01)
Q99 = df['Income'].quantile(0.99)

# cap outliers
df['Income_capped'] = df['Income'].clip(lower=Q1, upper=Q99)

# replace original column
df['Income'] = df['Income_capped']
df.drop(columns=['Income_capped'], inplace=True)

df['Income'].describe()
```

	count	2240.000000
min	51793.099444	
std	20616.153112	
min	7705.020000	
25%	35538.750000	
50%	51342.000000	
75%	68289.750000	
max	94437.680000	
Name:	Income, dtype: float64	

Step 6: Remove Unrealistic Age Values

- * Checked for customers with Age > 100
 - * Removed those rows from the dataset
 - * Verified that no unrealistic ages remain

Step 7: Encoding Categorical Variables

- * Applied ordinal encoding to Education
 - * Created one-hot encoded variables for Marital_Status
 - * Created one-hot encoded variables for Country
 - * Joined encoded columns back to the main DataFrame

```
[42]: # Ordinal encoding for Education:  
  
edu_order = {  
    'Basic': 1,  
    '2n_Cycle': 2,  
    'Graduation': 3,  
    'Master': 4,  
    'PhD': 5  
}  
  
# fix column name  
df['Education'] = df['Education'].str.replace('2n Cycle', '2n_Cycle')  
  
df['Education_ord'] = df['Education'].map(edu_order)  
  
df[['Education', 'Education_ord']].head()
```

```
[43]: # One-hot encoding for Marital_Status:  
  
marital_dummies = pd.get_dummies(df['Marital_Status'], prefix='Marital')  
df = df.concat([df, marital_dummies], axis=1)  
  
marital_dummies.head()
```

[43]:	Marital_Married	Marital_Previosuly_Married	Marital_Single
0	False	True	False
1	False	False	True
2	True	False	False
3	True	False	False
4	False	False	True

```
[44]: # One-hot encoding for Country:  
  
country_dummies = pd.get_dummies(df['Country'], prefix='Country')  
df = pd.concat([df, country_dummies], axis=1)  
  
country_dummies.head()  
  
[44]:   Country_AUS  Country_CA  Country_GER  Country_IND  Country_ME  Country_SA  Country_SP  Country_US  
0      False       False      False      False      False      False      True     False  
1      False       True       False      False      False      False      False     False  
2      False      False      False      False      False      False      False     True  
3      True       False      False      False      False      False      False     False  
4      False      False      False      False      False      False      True     False
```

Step 8: Correlation Heatmap

- * Select only numeric columns
- * Compute the correlation matrix
- * Plot a heatmap for visual interpretation

```
[45]: # Select numeric columns and compute correlation:
numeric_df = df.select_dtypes(include=['int64', 'float64'])
corr_matrix = numeric_df.corr()
corr_matrix
```

	ID	Year_Birth	Income	Kidhome	Teenhome	Recency	MntWines	MntFruits	MntMeatProducts	MntFishProducts	...	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	Response	C
ID	1.000000	0.003024	0.002879	0.002202	-0.003543	-0.046755	-0.021181	0.007080	-0.002622	-0.023181	...	-0.005062	-0.021524	-0.015027	-0.021810	
Year_Birth	0.003024	1.000000	-0.208115	0.234133	-0.363350	-0.019670	-0.163035	-0.013751	-0.030927	-0.042519	...	0.015322	-0.008227	-0.007657	0.018424	-
Income	0.002879	-0.208115	1.000000	-0.524802	0.040044	0.006729	0.7717841	0.528080	0.685489	0.542402	...	0.408476	0.336654	0.107633	0.168745	-
Kidhome	0.002202	0.234133	-0.524802	1.000000	-0.035753	0.007544	-0.496367	-0.372488	-0.437059	-0.387536	...	-0.204994	-0.172512	-0.081794	-0.080176	-
Teenhome	-0.003543	-0.363350	0.040044	-0.035753	1.000000	0.017115	0.005409	-0.175951	-0.260820	-0.203900	...	-0.190227	-0.140288	-0.015664	-0.154730	-
Recency	-0.046755	-0.019670	0.006729	0.007544	0.017115	1.000000	0.016668	-0.003592	0.023705	0.001532	...	0.000956	-0.019258	-0.001764	-0.198568	-
MntWines	-0.021181	-0.163035	0.717841	-0.496367	0.005409	0.016668	1.000000	0.388518	0.561993	0.399073	...	0.471969	0.354365	0.206040	0.247392	-
MntFruits	0.007080	-0.013751	0.528080	-0.372488	-0.175951	-0.003592	0.388518	1.000000	0.542057	0.594438	...	0.212027	0.195380	-0.009701	0.125904	-
MntMeatProducts	-0.002622	-0.030927	0.685489	-0.437059	-0.260820	0.023705	0.561993	0.542057	1.000000	0.567880	...	0.372212	0.310096	0.043090	0.236640	-
MntFishProducts	-0.023181	-0.042519	0.542402	-0.387536	-0.203900	0.001532	0.399073	0.594438	0.567880	1.000000	...	0.198163	0.260908	0.002583	0.111415	-
MntSweetProducts	-0.006444	-0.019571	0.538464	-0.370656	-0.162218	0.023045	0.385992	0.567054	0.523418	0.579535	...	0.258648	0.241875	0.009972	0.117366	-
MntGoldProd	-0.010661	-0.057599	0.404110	-0.349633	-0.020186	0.017412	0.386376	0.390042	0.348845	0.422103	...	0.176382	0.167145	0.050252	0.140693	-
NumDealsPurchases	-0.036917	-0.067999	-0.116465	0.221799	0.387792	-0.000987	0.010829	-0.131886	-0.122465	-0.139440	...	-0.182910	-0.123530	-0.037814	0.001854	-
NumWebPurchases	-0.017913	-0.153973	0.476716	-0.362063	0.155776	-0.010616	0.542177	0.297024	0.293579	0.293489	...	0.138958	0.154991	0.034103	0.148453	-
NumCatalogPurchases	-0.001893	-0.125439	0.690528	-0.502438	-0.110285	0.025449	0.634784	0.487307	0.723519	0.534033	...	0.321419	0.308240	0.099891	0.220894	-
NumStorePurchases	-0.014062	-0.139465	0.665315	-0.500387	0.050517	0.001117	0.642433	0.463168	0.480110	0.460099	...	0.216147	0.183043	0.085098	0.038855	-
NumWebVisitsMonth	-0.008104	0.117570	-0.646763	0.447641	0.134491	-0.021959	-0.320337	-0.417427	-0.539203	-0.445760	...	-0.276371	-0.192948	-0.007330	-0.004449	-
AcceptedCmp3	-0.035959	0.061013	-0.012898	0.014603	-0.042823	-0.032976	0.062201	0.014983	0.018331	0.000370	...	0.080930	0.094661	0.071981	0.254144	-
AcceptedCmp4	-0.025292	-0.064341	0.225366	-0.161775	0.038790	0.018890	0.373532	0.010402	0.103053	0.016864	...	0.307812	0.251225	0.292184	0.176690	-
AcceptedCmp5	-0.005062	0.015322	0.408476	-0.204994	-0.190227	0.000956	0.471969	0.212027	0.372212	0.198163	...	1.000000	0.404616	0.222333	0.328182	-

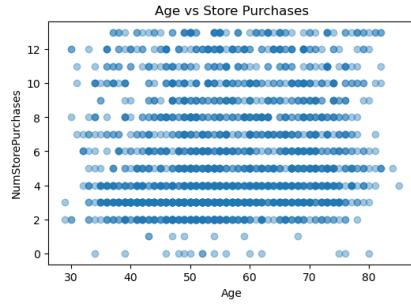


Step 9: Hypothesis Testing

9.1 - Hypothesis A - Age vs Store Purchases

- * Compared Age with number of in-store purchases
- * Visualized relationship using a scatter plot
- * Calculated Pearson correlation coefficient
- * Evaluated statistical significance

```
[48]: # Scatter plot:  
plt.figure(figsize=(6,4))  
plt.scatter(df['Age'], df['NumStorePurchases'], alpha=0.4)  
plt.xlabel('Age')  
plt.ylabel('NumStorePurchases')  
plt.title('Age vs Store Purchases')  
plt.show()
```



```
[49]: # Pearson correlation test:  
r, p = pearsonr(df['Age'], df['NumStorePurchases'])  
  
print("Correlation coefficient (r):", r)  
print("p-value:", p)  
  
Correlation coefficient (r): 0.13946510380279412  
p-value: 3.472816546739386e-11
```

9.2 - Hypothesis B - Children vs Online Purchases

- * Compared TotalChildren with number of online purchases
- * Visualized distributions using a boxplot
- * Performed the Spearman rank correlation test
- * Checked if online purchases increase with more children

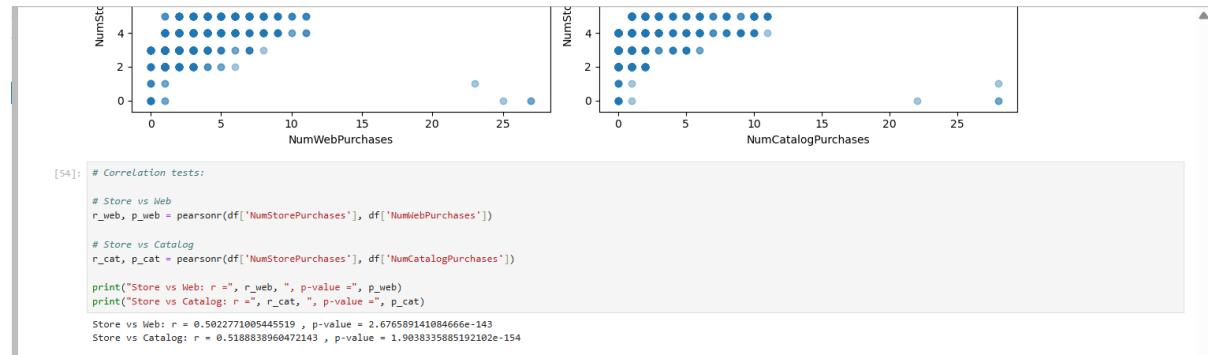
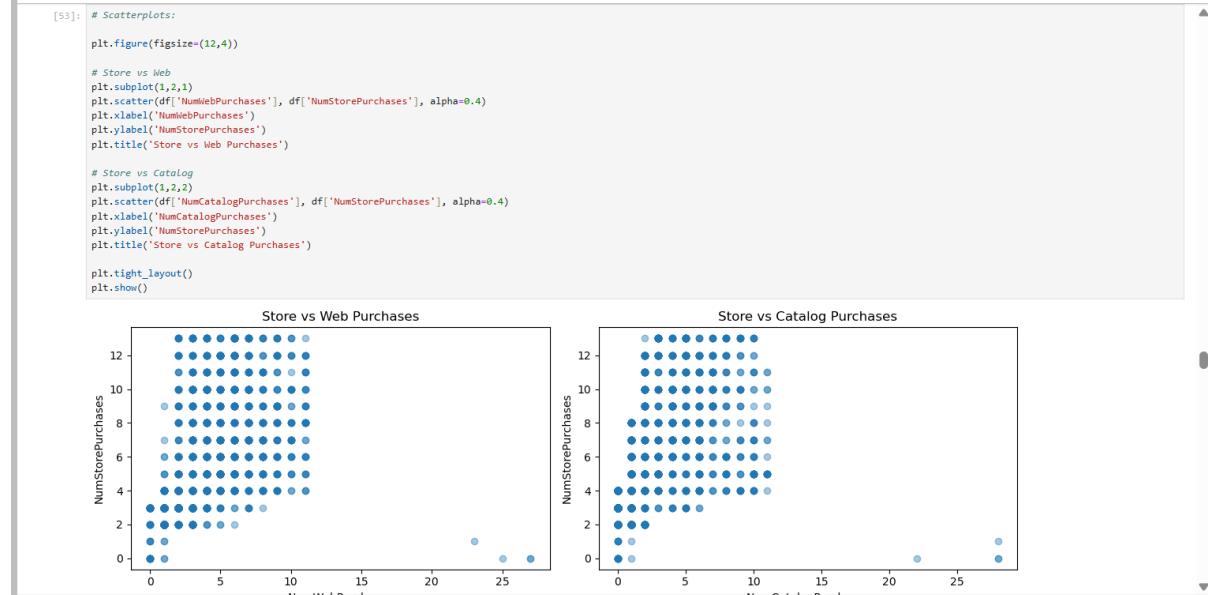
```
[50]: # Boxplot:  
plt.figure(figsize=(6,4))  
sns.boxplot(x=df['TotalChildren'], y=df['NumWebPurchases'])  
plt.xlabel('Total Children')  
plt.ylabel('NumWebPurchases')  
plt.title('Online Purchases by Number of Children')  
plt.show()
```



```
[52]: # Spearman correlation test:  
rho, p = spearmanr(df['TotalChildren'], df['NumWebPurchases'])  
  
print("Spearman rho:", rho)  
print("p-value:", p)  
  
Spearman rho: -0.185333569053294  
p-value: 9.791403824731e-19
```

9.3 - Hypothesis C — Store vs Web & Catalog Purchases

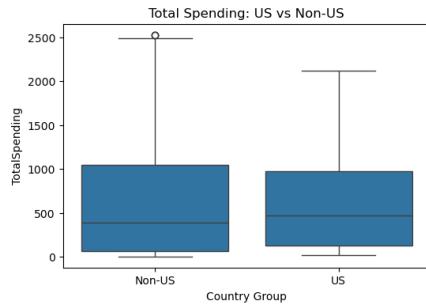
- * Visualized relationships using scatterplots
- * Performed Pearson correlation tests
- * Checked if online/catalog sales negatively impact store sales



9.4 - Hypothesis D - US vs Non-US Spending

- * Created two groups (US vs Non-US)
- * Visualized spending distributions
- * Performed Mann–Whitney U test
- * Checked if the difference is statistically significant

```
[55]: # Boxplot:  
plt.figure(figsize=(6,4))  
  
sns.boxplot(  
    x = df['Country'].apply(lambda x: 'US' if x=='US' else 'Non-US'),  
    y = df['TotalSpending'])  
plt.xlabel('Country Group')  
plt.ylabel('TotalSpending')  
plt.title('Total Spending: US vs Non-US')  
plt.show()
```

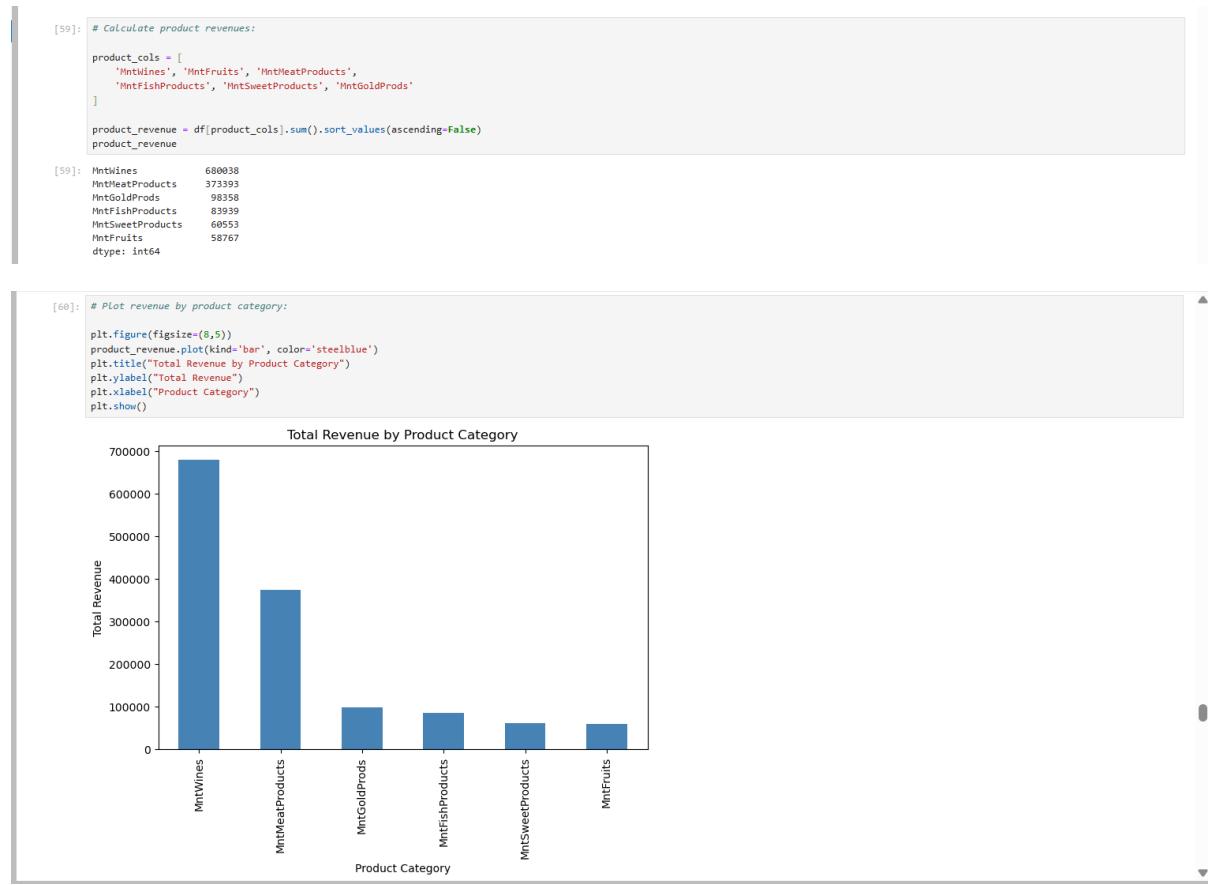


```
[56]: # Median spending:  
df.groupby(df['Country']).apply(lambda x: 'US' if x=='US' else 'Non-US'))['TotalSpending'].median()  
  
[56]: Country  
Non-US    393.0  
US        467.0  
Name: TotalSpending, dtype: float64  
  
[58]: # Mann-Whitney U test:  
us_spend = df[df['Country']=='US']['TotalSpending']  
nonus_spend = df[df['Country']!='US']['TotalSpending']  
  
stat, p = mannwhitneyu(us_spend, nonus_spend, alternative='two-sided')  
print("Mann-Whitney U statistic:", stat)  
print("p-value:", p)  
print("\nUS median:", us_spend.median())  
print("Non-US median:", nonus_spend.median())  
  
Mann-Whitney U statistic: 123822.5  
p-value: 0.23290054549364958  
US median: 467.0  
Non-US median: 393.0
```

Step 10: Visual Analysis

10.1 - Product Revenue Analysis

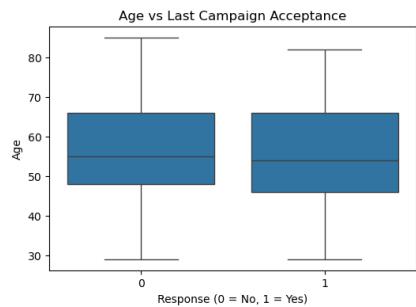
- * Calculated revenue across all product categories
- * Identified top and bottom performers
- * Visualized using a bar chart



10.2 - Age vs Campaign Acceptance

- * Compared Age distribution between responders and non-responders
- * Visualized using a boxplot
- * Calculated point-biserial correlation to measure association

```
[61]: # Boxplot of Age vs Response:  
plt.figure(figsize=(6,4))  
sns.boxplot(x=df['Response'], y=df['Age'])  
plt.xlabel("Response (0 = No, 1 = Yes)")  
plt.ylabel("Age")  
plt.title("Age vs Last Campaign Acceptance")  
plt.show()
```



```
[62]: # Median Age by Response:  
df.groupby('Response')['Age'].median()  
  
[62]: Response  
0    55.0  
1    54.0  
Name: Age, dtype: float64  
  
[64]: # Point-Biserial correlation:  
r, p = pointbiserialr(df['Response'], df['Age'])  
print("Point-biserial correlation (r):", r)  
print("p-value:", p)  
Point-biserial correlation (r): -0.018424292565464954  
p-value: 0.3837544666678645
```

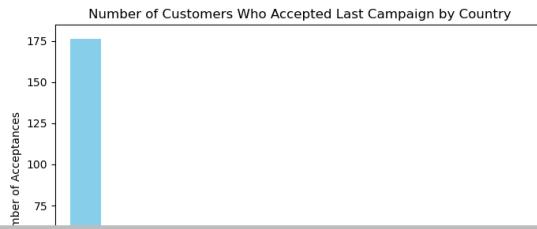
10.3 - Campaign Acceptance by Country

- * Grouped customers by country
- * Counted total number of acceptances
- * Visualized acceptance distribution by country

```
[65]: # Calculate acceptances by country:  
acceptance_by_country = df.groupby('Country')['Response'].sum().sort_values(ascending=False)
```

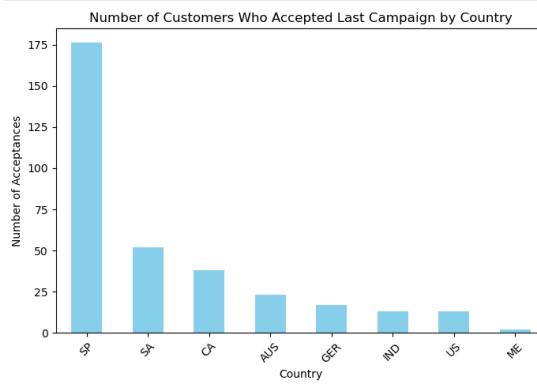
```
[65]:  
Country  
SP    176  
SA     52  
CA     38  
AUS    23  
GER    17  
IND    13  
US     13  
ME     2  
Name: Response, dtype: int64
```

```
[66]: # Plot acceptances by country:  
plt.figure(figsize=(8,5))  
acceptance_by_country.plot(kind='bar', color='skyblue')  
plt.title("Number of Customers Who Accepted Last Campaign by Country")  
plt.ylabel("Number of Acceptances")  
plt.xlabel("Country")  
plt.xticks(rotation=45)  
plt.show()
```



```
[66]: # Plot acceptances by country:
```

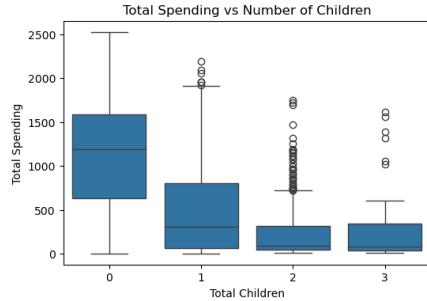
```
plt.figure(figsize=(8,5))  
acceptance_by_country.plot(kind='bar', color='skyblue')  
plt.title("Number of Customers Who Accepted Last Campaign by Country")  
plt.ylabel("Number of Acceptances")  
plt.xlabel("Country")  
plt.xticks(rotation=45)  
plt.show()
```



10.4 - Total Children vs Total Spending

- * Visualized spending for customers with different numbers of children
- * Calculated median spending across groups
- * Measured correlation strength using Spearman rank correlation

```
[70]: # Boxplot:  
plt.figure(figsize=(6,4))  
sns.boxplot(x=df['TotalChildren'], y=df['TotalSpending'])  
plt.xlabel('Total Children')  
plt.ylabel('Total Spending')  
plt.title('Total Spending vs Number of Children')  
plt.show()
```



```
[71]: # Median spending by children group:  
df.groupby('TotalChildren')['TotalSpending'].median()
```

```
[71]: TotalChildren  
0    1189.0  
1     306.0  
2      93.0  
3      88.0  
Name: TotalSpending, dtype: float64
```

```
[72]: # Spearman correlation test:  
rho, p = spearmanr(df['TotalChildren'], df['TotalSpending'])  
print("Spearman rho:", rho)  
print("p-value:", p)  
Spearman rho: -0.4835857068087243  
p-value: 1.817901317197658e-131
```

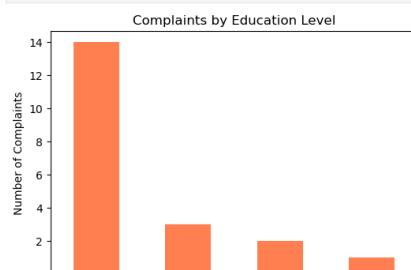
10.5 - Complaints by Education Level

- * Filtered only customers who complained
- * Counted complaints by education category
- * Visualized distribution using a bar chart

```
[73]: # Count complaints by education:  
complaints_by_edu = df[df['Complain'] == 1]['Education'].value_counts()  
complaints_by_edu
```

```
[73]: Education  
Graduation    14  
2n_Cycle      3  
Master        2  
PhD           1  
Name: count, dtype: int64
```

```
[74]: # Plot complaints by education:  
plt.figure(figsize=(6,4))  
complaints_by_edu.plot(kind='bar', color='coral')  
plt.title("Complaints by Education Level")  
plt.xlabel("Education Level")  
plt.ylabel("Number of Complaints")  
plt.xticks(rotation=45)  
plt.show()
```



Step 11: Final Project Summary

Final Summary — Marketing Campaign Analysis

- * Completed full data cleaning and preprocessing
- * Engineered new features for deeper insights
- * Performed exploratory data analysis
- * Conducted outlier treatment and encoding
- * Executed four hypothesis tests as required
- * Generated key visual insights for business interpretation

Key data preparation steps:

- * Cleaned Income, Education, and Marital_Status categories
- * Converted Dt_Customer to datetime format
- * Imputed missing income values using Education + Marital_Status groups
- * Removed unrealistic Age values
- * Created TotalChildren, Age, TotalSpending, and TotalPurchases

Key EDA findings:

- * Income and spending variables are right-skewed
- * Wines and Meat are the highest revenue products
- * TotalSpending strongly correlates with premium product categories
- * Multi-channel shoppers tend to spend more

Hypothesis testing results:

- * Older customers show very weak preference for in-store shopping
- * Customers with more children do not prefer online shopping; they shop less
- * No evidence of channel cannibalization — store, web, and catalog purchases rise together
- * US customers do not significantly outperform non-US customers in spending

Key visual insights:

- * Spain has the highest campaign acceptance
- * Campaign acceptance is not influenced by Age
- * Families with more children spend significantly less
- * Most complaints come from Graduation-level customers

Business insights:

- * Spain is highly responsive — stronger marketing focus recommended
- * High-value customers tend to have no children
- * Multi-channel engagement drives higher spending
- * Graduation-level customers may need better support or communication

